

Subgradient Methods

Stephen Boyd (with help from Jaehyun Park)
Notes for EE364b, Stanford University, Spring 2013–14

May 2014; based on notes from January 2007

Contents

1	Introduction	3
2	Basic subgradient method	4
2.1	Negative subgradient update	4
2.2	Step size rules	4
2.3	Convergence results	5
3	Convergence proof	6
3.1	Assumptions	6
3.2	Some basic inequalities	6
3.3	A bound on the suboptimality bound	8
3.4	A stopping criterion	9
3.5	Numerical example	9
4	Polyak’s step length	12
4.1	Optimal step size choice when f^* is known	12
4.2	Polyak step size choice with estimated f^*	12
4.3	Numerical example	13
5	Alternating projections	16
5.1	Finding a point in the intersection of convex sets	16
5.2	Solving convex inequalities	17
5.3	Positive semidefinite matrix completion	18
6	Projected subgradient method	21
6.1	Numerical example	22
6.2	Projected subgradient for dual problem	23
6.3	Numerical example	24

7	Subgradient method for constrained optimization	26
7.1	Numerical example	28
8	Primal-dual subgradient method	30
8.1	Equality constrained problems	30
8.2	Inequality constrained problems	33
8.3	Numerical example	35
9	Speeding up subgradient methods	37
9.1	Heavy ball methods	37

1 Introduction

The *subgradient method* is a very simple algorithm for minimizing a nondifferentiable convex function. The method looks very much like the ordinary gradient method for differentiable functions, but with several notable exceptions:

- The subgradient method applies directly to *nondifferentiable* f .
- The step lengths are not chosen via a line search, as in the ordinary gradient method. In the most common cases, the step lengths are fixed ahead of time.
- Unlike the ordinary gradient method, the subgradient method is *not* a descent method; the function value can (and often does) increase.

The subgradient method is readily extended to handle problems with constraints.

Subgradient methods can be *much* slower than interior-point methods (or Newton's method in the unconstrained case). In particular, they are first-order methods; their performance depends very much on the problem scaling and conditioning. (In contrast, Newton and interior-point methods are second-order methods, not affected by problem scaling.)

However, subgradient methods do have some advantages over interior-point and Newton methods. They can be immediately applied to a far wider variety of problems than interior-point or Newton methods. The memory requirement of subgradient methods can be much smaller than an interior-point or Newton method, which means it can be used for extremely large problems for which interior-point or Newton methods cannot be used. Moreover, by combining the subgradient method with primal or dual decomposition techniques, it is sometimes possible to develop a simple distributed algorithm for a problem. In any case, subgradient methods are well worth knowing about.

The subgradient method was originally developed by Shor and others in the Soviet Union in the 1960s and 1970s. A basic reference on subgradient methods is his book [Sho85]; a very clear discussion can be found in chapter 5 of Polyak's book [Pol87]. Bertsekas [Ber99] is another good reference on the subgradient method, in particular, on how to combine it with primal and dual decomposition. Other book treatments of the topic are in Ruszczyński [Rus06, §7.1], Nesterov [Nes04, Chap. 3], Akgul [Akg84], Yudin and Nemirovski [NY83], Censor and Zenios [CZ97], and Shor [Sho98, Chap. 2]. Some interesting recent research papers on subgradient methods are [NB01] and [Nes09].

2 Basic subgradient method

2.1 Negative subgradient update

We start with the unconstrained case, where the goal is to minimize $f : \mathbf{R}^n \rightarrow \mathbf{R}$, which is convex and has domain \mathbf{R}^n (for now). To do this, the subgradient method uses the simple iteration

$$x^{(k+1)} = x^{(k)} - \alpha_k g^{(k)}.$$

Here $x^{(k)}$ is the k th iterate, $g^{(k)}$ is any subgradient of f at $x^{(k)}$, and $\alpha_k > 0$ is the k th step size. Thus, at each iteration of the subgradient method, we take a step in the direction of a negative subgradient.

Recall that a subgradient of f at x is any vector g that satisfies the inequality $f(y) \geq f(x) + g^T(y - x)$ for all y . When f is differentiable, the only possible choice for $g^{(k)}$ is $\nabla f(x^{(k)})$, and the subgradient method then reduces to the gradient method (except, as we'll see below, for the choice of step size). The set of subgradients of f at x is the subdifferential of f at x , denoted $\partial f(x)$. So the condition that $g^{(k)}$ be a subgradient of f at $x^{(k)}$ can be written $g^{(k)} \in \partial f(x^{(k)})$.

~~It can happen that~~ $-g^{(k)}$ is not a descent direction for f at $x^{(k)}$, i.e., $f'(x; -g^{(k)}) > 0$. In such cases we always have $f(x^{(k+1)}) > f(x^{(k)})$. Even when $-g^{(k)}$ is a descent direction at $x^{(k)}$, the step size can be such $f(x^{(k+1)}) > f(x^{(k)})$. In other words, an iteration of the subgradient method can increase the objective function.

Since the subgradient method is not a descent method, it is common to keep track of the best point found so far, i.e., the one with smallest function value. At each step, we set

$$f_{\text{best}}^{(k)} = \min\{f_{\text{best}}^{(k-1)}, f(x^{(k)})\},$$

and set $i_{\text{best}}^{(k)} = k$ if $f(x^{(k)}) = f_{\text{best}}^{(k)}$, i.e., if $x^{(k)}$ is the best point found so far. (In a descent method there is no need to do this, since the current point is always the best one so far.) Then we have

$$f_{\text{best}}^{(k)} = \min\{f(x^{(1)}), \dots, f(x^{(k)})\},$$

i.e., the best objective value found in k iterations. Since $f_{\text{best}}^{(k)}$ is decreasing, it has a limit (which can be $-\infty$).

2.2 Step size rules

In the subgradient method the step size selection is very different from the standard gradient method. Many different types of step size rules are used. We'll start with five basic step size rules.

- *Constant step size.* $\alpha_k = \alpha$ is a positive constant, independent of k .
- *Constant step length.* $\alpha_k = \gamma / \|g^{(k)}\|_2$, where $\gamma > 0$. This means that $\|x^{(k+1)} - x^{(k)}\|_2 = \gamma$.

- *Square summable but not summable.* The step sizes satisfy

$$\alpha_k \geq 0, \quad \sum_{k=1}^{\infty} \alpha_k^2 < \infty, \quad \sum_{k=1}^{\infty} \alpha_k = \infty.$$

One typical example is $\alpha_k = a/(b+k)$, where $a > 0$ and $b \geq 0$.

- *Nonsummable diminishing.* The step sizes satisfy

$$\alpha_k \geq 0, \quad \lim_{k \rightarrow \infty} \alpha_k = 0, \quad \sum_{k=1}^{\infty} \alpha_k = \infty.$$

Step sizes that satisfy this condition are called *diminishing step size rules*. A typical example is $\alpha_k = a/\sqrt{k}$, where $a > 0$.

- *Nonsummable diminishing step lengths.* The step sizes are chosen as $\alpha_k = \gamma_k / \|g^{(k)}\|_2$, where

$$\gamma_k \geq 0, \quad \lim_{k \rightarrow \infty} \gamma_k = 0, \quad \sum_{k=1}^{\infty} \gamma_k = \infty.$$

There are still other choices, and many variations on these choices. In §4.1 we will encounter another step size rule that requires knowledge of the optimal value f^* .

The most interesting feature of these choices is that *they are determined before the algorithm is run; they do not depend on any data computed during the algorithm*. This is very different from the step size rules found in standard descent methods, which very much depend on the current point and search direction.

2.3 Convergence results

There are many results on convergence of the subgradient method. For constant step size and constant step length, the subgradient algorithm is guaranteed to converge to within some range of the optimal value, i.e., we have

$$\lim_{k \rightarrow \infty} f_{\text{best}}^{(k)} - f^* < \epsilon,$$

where f^* denotes the optimal value of the problem, i.e., $f^* = \inf_x f(x)$. (This implies that the subgradient method finds an ϵ -suboptimal point within a finite number of steps.) The number ϵ is a function of the step size parameter h , and decreases with it.

For the diminishing step size and step length rules (and therefore also the square summable but not summable step size rule), the algorithm is guaranteed to converge to the optimal value, i.e., we have $\lim_{k \rightarrow \infty} f(x^{(k)}) = f^*$. It's remarkable that such a simple algorithm can be used to minimize any convex function for which you can compute a subgradient at each point. We'll also see that the convergence proof is also simple.

When the function f is differentiable, we can say a bit more about the convergence. In this case, the subgradient method with constant step size yields convergence to the optimal value, provided the parameter α is small enough.

3 Convergence proof

3.1 Assumptions

Here we give a proof of some typical convergence results for the subgradient method. We assume that there is a minimizer of f , say x^* . We also make one other assumption on f : We will assume that the norm of the subgradients is bounded, *i.e.*, there is a G such that $\|g^{(k)}\|_2 \leq G$ for all k . This will be the case if, for example, f satisfies the Lipschitz condition

$$|f(u) - f(v)| \leq G\|u - v\|_2,$$

for all u, v , because then $\|g\|_2 \leq G$ for any $g \in \partial f(x)$, and any x . In fact, some versions of the subgradient method (*e.g.*, diminishing nonsummable step lengths) work when this assumption doesn't hold; see [Sho85] or [Pol87].

We'll also assume that a number R is known that satisfies $R \geq \|x^{(1)} - x^*\|_2$. We can interpret R as an upper bound on $\mathbf{dist}(x^{(1)}, X^*)$, the distance of the initial point to the optimal set.

3.2 Some basic inequalities

For the standard gradient descent method, the convergence proof is based on the function value decreasing at each step. In the subgradient method, the key quantity is not the function value (which often increases); it is the *Euclidean distance to the optimal set*.

Recall that x^* is a point that minimizes f , *i.e.*, it is an arbitrary optimal point. We have

$$\begin{aligned} \|x^{(k+1)} - x^*\|_2^2 &= \|x^{(k)} - \alpha_k g^{(k)} - x^*\|_2^2 \\ &= \|x^{(k)} - x^*\|_2^2 - 2\alpha_k g^{(k)T}(x^{(k)} - x^*) + \alpha_k^2 \|g^{(k)}\|_2^2 \\ &\leq \|x^{(k)} - x^*\|_2^2 - 2\alpha_k (f(x^{(k)}) - f^*) + \alpha_k^2 \|g^{(k)}\|_2^2, \end{aligned}$$

where $f^* = f(x^*)$. The last line follows from the definition of subgradient, which gives

$$f(x^*) \geq f(x^{(k)}) + g^{(k)T}(x^* - x^{(k)}).$$

Applying the inequality above recursively, we have

$$\|x^{(k+1)} - x^*\|_2^2 \leq \|x^{(1)} - x^*\|_2^2 - 2 \sum_{i=1}^k \alpha_i (f(x^{(i)}) - f^*) + \sum_{i=1}^k \alpha_i^2 \|g^{(i)}\|_2^2.$$

Using $\|x^{(k+1)} - x^*\|_2^2 \geq 0$ and $\|x^{(1)} - x^*\|_2 \leq R$ we have

$$\underline{2 \sum_{i=1}^k \alpha_i (f(x^{(i)}) - f^*) \leq R^2 + \sum_{i=1}^k \alpha_i^2 \|g^{(i)}\|_2^2.} \tag{1}$$

Combining this with

$$\sum_{i=1}^k \alpha_i (f(x^{(i)}) - f^*) \geq \left(\sum_{i=1}^k \alpha_i \right) \min_{i=1, \dots, k} (f(x^{(i)}) - f^*) = \left(\sum_{i=1}^k \alpha_i \right) (f_{\text{best}}^{(k)} - f^*),$$

we have the inequality

$$f_{\text{best}}^{(k)} - f^* = \min_{i=1,\dots,k} f(x^{(i)}) - f^* \leq \frac{R^2 + \sum_{i=1}^k \alpha_i^2 \|g^{(i)}\|_2^2}{2 \sum_{i=1}^k \alpha_i}. \quad (2)$$

Finally, using the assumption $\|g^{(k)}\|_2 \leq G$, we obtain the basic inequality

$$\underline{f_{\text{best}}^{(k)} - f^*} \leq \frac{R^2 + G^2 \sum_{i=1}^k \alpha_i^2}{2 \sum_{i=1}^k \alpha_i}. \quad (3)$$

From this inequality we can read off various convergence results.

Constant step size. When $\alpha_k = \alpha$, we have

$$f_{\text{best}}^{(k)} - f^* \leq \frac{R^2 + G^2 \alpha^2 k}{2 \alpha k}.$$

The righthand side converges to $G^2 \alpha / 2$ as $k \rightarrow \infty$. Thus, for the subgradient method with fixed step size α , $f_{\text{best}}^{(k)}$ converges to within $G^2 \alpha / 2$ of optimal. We also find that $f(x^{(k)}) - f^* \leq G^2 \alpha$ within at most $R^2 / (G^2 \alpha^2)$ steps.

Constant step length. With $\alpha_k = \gamma / \|g^{(k)}\|_2$, the inequality (2) becomes

$$f_{\text{best}}^{(k)} - f^* \leq \frac{R^2 + \gamma^2 k}{2 \sum_{i=1}^k \alpha_i} \leq \frac{R^2 + \gamma^2 k}{2 \gamma k / G},$$

using $\alpha_i \geq \gamma / G$. The righthand side converges to $G \gamma / 2$ as $k \rightarrow \infty$, so in this case the subgradient method converges to within $G \gamma / 2$ of optimal.

Square summable but not summable. Now suppose

$$\|\alpha\|_2^2 = \sum_{k=1}^{\infty} \alpha_k^2 < \infty, \quad \sum_{k=1}^{\infty} \alpha_k = \infty.$$

Then we have

$$f_{\text{best}}^{(k)} - f^* \leq \frac{R^2 + G^2 \|\alpha\|_2^2}{2 \sum_{i=1}^k \alpha_i},$$

which converges to zero as $k \rightarrow \infty$, since the numerator converges to $R^2 + G^2 \|\alpha\|_2^2$, and the denominator grows without bound. Thus, the subgradient method converges (in the sense $f_{\text{best}}^{(k)} \rightarrow f^*$).

Diminishing step size rule. If the sequence α_k converges to zero and is nonsummable, then the righthand side of the inequality (3) converges to zero, which implies the subgradient method converges. To show this, let $\epsilon > 0$. Then there exists an integer N_1 such that $\alpha_i \leq \epsilon/G^2$ for all $i > N_1$. There also exists an integer N_2 such that

$$\sum_{i=1}^{N_2} \alpha_i \geq \frac{1}{\epsilon} \left(R^2 + G^2 \sum_{i=1}^{N_1} \alpha_i^2 \right),$$

since $\sum_{i=1}^{\infty} \alpha_i = \infty$. Let $N = \max\{N_1, N_2\}$. Then for $k > N$, we have

$$\begin{aligned} \frac{R^2 + G^2 \sum_{i=1}^k \alpha_i^2}{2 \sum_{i=1}^k \alpha_i} &\leq \frac{R^2 + G^2 \sum_{i=1}^{N_1} \alpha_i^2}{2 \sum_{i=1}^k \alpha_i} + \frac{G^2 \sum_{i=N_1+1}^k \alpha_i^2}{2 \sum_{i=1}^{N_1} \alpha_i + 2 \sum_{i=N_1+1}^k \alpha_i} \\ &\leq \frac{R^2 + G^2 \sum_{i=1}^{N_1} \alpha_i^2}{(2/\epsilon) \left(R^2 + G^2 \sum_{i=1}^{N_1} \alpha_i^2 \right)} + \frac{G^2 \sum_{i=N_1+1}^k (\epsilon \alpha_i / G^2)}{2 \sum_{i=N_1+1}^k \alpha_i} \\ &= \frac{\epsilon}{2} + \frac{\epsilon}{2} = \epsilon. \end{aligned}$$

Nonsummable diminishing step lengths. Finally, suppose that $\alpha_k = \gamma_k / \|g^{(k)}\|_2$, with γ_k nonsummable and converging to zero. The inequality (2) becomes

$$f_{\text{best}}^{(k)} - f^* \leq \frac{R^2 + \sum_{i=1}^k \gamma_i^2}{2 \sum_{i=1}^k \alpha_i} \leq \frac{R^2 + \sum_{i=1}^k \gamma_i^2}{(2/G) \sum_{i=1}^k \gamma_i},$$

which converges to zero as $k \rightarrow \infty$.

3.3 A bound on the suboptimality bound

It's interesting to ask the question, what sequence of step sizes minimizes the righthand side of (3)? In other words, how do we choose positive $\alpha_1, \dots, \alpha_k$ so that

$$\frac{R^2 + G^2 \sum_{i=1}^k \alpha_i^2}{2 \sum_{i=1}^k \alpha_i}$$

(which is an upper bound on $f_{\text{best}}^{(k)} - f^*$) is minimized? This is a convex and symmetric function of $\alpha_1, \dots, \alpha_k$, so we conclude the optimal occurs when all α_i are equal (to, say, α). This reduces our suboptimality bound to

$$\frac{R^2 + G^2 k \alpha^2}{2k\alpha}$$

which is minimized by $\alpha = (R/G)/\sqrt{k}$.

In other words, the choice of $\alpha_1, \dots, \alpha_k$ that minimizes the suboptimality bound (3) is given by

$$\alpha_i = (R/G)/\sqrt{k}, \quad i = 1, \dots, k.$$

This choice of constant step size yields the suboptimality bound

$$f_{\text{best}}^{(k)} - f^* \leq RG/\sqrt{k}.$$

Put another way, we can say that for *any* choice of step sizes, the suboptimality bound (3) must be at least as large as RG/\sqrt{k} . If we use (3) as our stopping criterion, then the number of steps to achieve a guaranteed accuracy of ϵ will be at least $(RG/\epsilon)^2$, no matter what step sizes we use. (It will be this number if we use the step size $\alpha_k = (R/G)/\sqrt{k}$).

Note that RG has a simple interpretation as an initial bound on $f(x^{(1)}) - f^*$, based on $\|x^{(1)} - x^*\|_2 \leq R$ and the Lipschitz constant G for f . Thus $(RG)/\epsilon$ is the ratio of initial uncertainty in f^* to final uncertainty in f^* . If we square this number, we get the minimum number of steps it will take to achieve this reduction in uncertainty. This tells us that the subgradient method is going to be very slow, if we use (3) as our stopping criterion. To reduce the initial uncertainty by a factor of 1000, say, it will require at least 10^6 iterations.

3.4 A stopping criterion

We can use (1) to find a lower bound on f^* that is sharper than the lower bounds (2) and (3), and can be used as a stopping criterion. Re-arranging (1) and using $R \geq \|x^{(1)} - x^*\|_2$, we get

$$f^* \geq l_k = \frac{2 \sum_{i=1}^k \alpha_i f(x^{(i)}) - R^2 - \sum_{i=1}^k \alpha_i^2 \|g^{(i)}\|_2^2}{2 \sum_{i=1}^k \alpha_i}, \quad (4)$$

which can be computed after the k th step. The sequence l_1, l_2, \dots need not increase, so we can keep track of the best lower bound on f^* found so far,

$$l_{\text{best}}^{(k)} = \max\{l_1, \dots, l_k\}.$$

We can terminate the algorithm when $f_{\text{best}}^{(k)} - l_{\text{best}}^{(k)}$ is smaller than some threshold.

This bound is better than (3), and doesn't depend on G , but it too goes to zero very slowly. For this reason, the subgradient method is usually used without any formal stopping criterion.

3.5 Numerical example

We consider the problem of minimizing a piecewise linear function:

$$\text{minimize } f(x) = \max_{i=1, \dots, m} (a_i^T x + b_i),$$

with variable $x \in \mathbf{R}^n$. Of course this problem is readily (and efficiently) solved via linear programming.

Finding a subgradient of f is easy: given x , we first find an index j for which

$$a_j^T x + b_j = \max_{i=1, \dots, m} (a_i^T x + b_i).$$




Figure 1: The value of $f_{\text{best}}^{(k)} - f^*$ versus iteration number k , for the subgradient method with constant step length γ .

Then we can take as subgradient $g = a_j$. We can take $G = \max_{i=1,\dots,m} \|a_i\|_2$.

We illustrate the subgradient method with a specific problem instance with $n = 20$ variables and $m = 100$ terms, with problem data a_i and b_i generated from a unit normal distribution. We start with $x^{(1)} = 0$. There is no simple way to find a justifiable value for R (*i.e.*, a value of R for which we can prove that $\|x^{(1)} - x^*\|_2 \leq R$ holds) so we take $R = 10$. For our particular problem instance, it turns out that $\|x^{(1)} - x^*\|_2 = 0.91$, where we computed an optimal point and the optimal value $f^* \approx 1.1$ using linear programming.

We first consider the constant step length rule $\alpha_k = \gamma / \|g^{(k)}\|_2$. Figure 1 shows convergence of $f_{\text{best}}^{(k)} - f^*$ for $\gamma = 0.05$, $\gamma = 0.01$, and $\gamma = 0.005$. The figure reveals a trade-off: larger γ gives faster convergence, but larger final suboptimality.

To illustrate the subgradient method with some diminishing step size rules, we consider the nonsummable diminishing step size rule $\alpha_k = 0.1/\sqrt{k}$, and the square summable but not summable step rule $\alpha_k = 1/k$. The convergence for these step size rules is plotted in figure 2.

These plots are fairly typical: The subgradient method is very slow. But what do you expect from an algorithm that is just a few lines of code, has no line search, and uses any subgradient? (And has a convergence proof that is also just a few lines long.) One of its advantages, apart from simplicity, is robustness. We'll see this very clearly when we study the stochastic subgradient method.




Figure 2: The value of $f_{\text{best}}^{(k)} - f^*$ versus iteration number k , for the subgradient method with two diminishing step rules $\alpha_k = 0.1/\sqrt{k}$ and $\alpha_k = 1/\sqrt{k}$, and with two square summable step size rules $\alpha_k = 1/k$ and $\alpha_k = 10/k$.

4 Polyak's step length

In this section we describe a subgradient step length choice due to Polyak.

4.1 Optimal step size choice when f^\star is known

Polyak [Pol87] suggests a step size that can be used when the optimal value f^\star is known, and is in some sense optimal. (You might imagine that f^\star is rarely known, but we will see that's not the case.) The step size is

$$\alpha_k = \frac{f(x^{(k)}) - f^\star}{\|g^{(k)}\|_2^2}. \quad (5)$$

To motivate this step size, imagine that

$$f(x^{(k)} - \alpha g^{(k)}) \approx f(x^{(k)}) + g^{(k)T} (x^{(k)} - \alpha g^{(k)} - x^{(k)}) = f(x^{(k)}) - \alpha g^{(k)T} g^{(k)}.$$

(This would be the case if α were small, and $g^{(k)} = \nabla f(x^{(k)})$.) Replacing the lefthand side with f^\star and solving for α gives the step length above.

We can give another simple motivation for the step length (5). The subgradient method starts from the basic inequality

$$\|x^{(k+1)} - x^\star\|_2^2 \leq \|x^{(k)} - x^\star\|_2^2 - 2\alpha_k(f(x^{(k)}) - f^\star) + \alpha_k^2 \|g^{(k)}\|_2^2.$$

The step size (5) minimizes the righthand side.

To analyze convergence, we substitute the step size (5) into (1), to get

$$2 \sum_{i=1}^k \frac{(f(x^{(i)}) - f^\star)^2}{\|g^{(i)}\|_2^2} \leq R^2 + \sum_{i=1}^k \frac{(f(x^{(i)}) - f^\star)^2}{\|g^{(i)}\|_2^2},$$

so

$$\sum_{i=1}^k \frac{(f(x^{(i)}) - f^\star)^2}{\|g^{(i)}\|_2^2} \leq R^2.$$

Using $\|g^{(i)}\|_2 \leq G$ we get

$$\sum_{i=1}^k (f(x^{(i)}) - f^\star)^2 \leq R^2 G^2.$$

We conclude that $f(x^{(k)}) \rightarrow f^\star$. The number of steps needed before we can guarantee suboptimality ϵ is $k = (RG/\epsilon)^2$, which is optimal from our analysis above.

4.2 Polyak step size choice with estimated f^\star

The basic idea is to estimate the optimal value f^\star , as $f_{\text{best}} - \gamma^k$, where $\gamma^k > 0$ and $\gamma^k \rightarrow 0$. This suggests the step size

$$\alpha_k = \frac{f(x^{(k)}) - f_{\text{best}} + \gamma_k}{\|g^{(k)}\|_2^2}.$$

We'll also need $\sum_{k=1}^{\infty} \gamma_k = \infty$. Note that γ_k has a simple interpretation: It's our estimate of how suboptimal the current point is. Then we have $f_{\text{best}}^{(k)} \rightarrow f^*$.

To show this, we substitute α_i into the basic inequality (1) to get

$$\begin{aligned} R^2 &\geq \sum_{i=1}^k \left(2\alpha_i(f(x^{(i)}) - f^*) - \alpha_i^2 \|g^{(i)}\|_2^2 \right) \\ &= \sum_{i=1}^k \frac{2(f(x^{(i)}) - f_{\text{best}}^{(i)} + \gamma_i)(f(x^{(i)}) - f^*) - (f(x^{(i)}) - f_{\text{best}}^{(i)} + \gamma_i)^2}{\|g^{(i)}\|_2^2} \\ &= \sum_{i=1}^k \frac{(f(x^{(i)}) - f_{\text{best}}^{(i)} + \gamma_i)((f(x^{(i)}) - f^*) + (f_{\text{best}}^{(i)} - f^*) - \gamma_i)}{\|g^{(i)}\|_2^2}. \end{aligned}$$

Now we can prove convergence. Suppose $f_{\text{best}}^{(k)} - f^* \geq \epsilon > 0$. Then for $i = 1, \dots, k$, $f(x^{(i)}) - f^* \geq \epsilon$. Find N for which $\gamma_i \leq \epsilon$ for $i \geq N$. This implies the second term in the numerator is at least ϵ :

$$(f(x^{(i)}) - f^*) + (f_{\text{best}}^{(i)} - f^*) - \gamma_i \geq \epsilon.$$

In particular, it is positive. It follows the terms in the sum above for $i \geq N$ are positive. Let S denote the sum above, up to $i = N - 1$. (We assume $k \geq N$.) We then have

$$\sum_{i=N}^k \frac{(f(x^{(i)}) - f_{\text{best}}^{(i)} + \gamma_i)((f(x^{(i)}) - f^*) + (f_{\text{best}}^{(i)} - f^*) - \gamma_i)}{\|g^{(i)}\|_2^2} \leq R^2 - S.$$

We get a lower bound on the lefthand side using

$$f(x^{(i)}) - f_{\text{best}}^{(i)} + \gamma_i \geq \gamma_i$$

along with the inequality above and $\|g^{(i)}\|_2^2 \leq G$ to get

$$(\epsilon/G^2) \sum_{i=N}^k \gamma_i \leq R^2 - S.$$

Since the lefthand side converges to ∞ and righthand side doesn't depend on k , we see that k cannot be too large.

4.3 Numerical example

Figure 3 shows the progress of the subgradient method with Polyak's step size for the piecewise linear example from §3.5. Of course this isn't fair, since we don't know f^* before solving the problem. But this plot shows that even with this unfair advantage in choosing step lengths, the subgradient method is pretty slow.

Figure 4 shows f_{best}^k for the same example, with both the optimal step size, and the estimated step size using $\gamma_k = 10/(10 + k)$. In this example the Polyak step size with estimated optimal value is just as good as the step size with known optimal value.




Figure 3: The value of $f_{\text{best}}^{(k)} - f^*$ versus iteration number k , for the subgradient method with Polyak's step size (solid black line) and the subgradient methods with diminishing step sizes considered in the previous example (dashed lines).




Figure 4: The value of $f_{\text{best}}^{(k)} - f^*$ versus iteration number k , for the subgradient method with Polyak's step size (solid black line) and the estimated optimal step size (dashed red line).

5 Alternating projections

Polyak's step length can be used to derive some versions of the alternating projections method for finding a point in the intersection of convex sets.

5.1 Finding a point in the intersection of convex sets

Suppose we want to find a point in

$$C = C_1 \cap \cdots \cap C_m,$$

where $C_1, \dots, C_m \subseteq \mathbf{R}^n$ are closed and convex, and we assume that C is nonempty. We can do this by minimizing the function

$$f(x) = \max\{\mathbf{dist}(x, C_1), \dots, \mathbf{dist}(x, C_m)\},$$

which is convex, and has minimum value $f^\star = 0$ (since C is nonempty).

We first explain how to find a subgradient g of f at x . If $f(x) = 0$, we can take $g = 0$ (which in any case means we are done). Otherwise find an index j such that $\mathbf{dist}(x, C_j) = f(x)$, *i.e.*, find a set that has maximum distance to x . A subgradient of f is

$$g = \nabla \mathbf{dist}(x, C_j) = \frac{x - \Pi_{C_j}(x)}{\|x - \Pi_{C_j}(x)\|_2},$$

where Π_{C_j} is Euclidean projection onto C_j . Note that $\|g\|_2 = 1$, so we can take $G = 1$.

The subgradient algorithm update, with step size rule (5), and assuming that the index j is one for which $x^{(k)}$ has maximum distance to C_j , is given by

$$\begin{aligned} x^{(k+1)} &= x^{(k)} - \alpha_k g^{(k)} \\ &= x^{(k)} - f(x^{(k)}) \frac{x^{(k)} - \Pi_{C_j}(x^{(k)})}{\|x^{(k)} - \Pi_{C_j}(x^{(k)})\|_2} \\ &= \Pi_{C_j}(x^{(k)}). \end{aligned}$$

Here we use $\|g^{(k)}\|_2 = 1$ and $f^\star = 0$ in the second line, and

$$f(x^{(k)}) = \mathbf{dist}(x^{(k)}, C_j) = \|x^{(k)} - \Pi_{C_j}(x^{(k)})\|_2$$

in the third line.

The algorithm is very simple: at each step, we simply project the current point onto the farthest set. This is an extension of the famous *alternating projections* algorithm. (When there are just two sets, then at each step you project the current point onto the *other* set. Thus the projections simply alternate.)

We are only guaranteed that $f(x^{(k)}) \rightarrow f^\star = 0$. In other words, a subsequence of our points approaches a point in C ; we are not guaranteed to actually find a point in C (except in the limit). This can be addressed several ways. One way is to run the algorithm using

closed sets $\tilde{C}_i \subseteq \text{int } C_i$, so that $x^{(k)} \rightarrow \tilde{C} = \tilde{C}_1 \cap \cdots \cap \tilde{C}_m$. Then we are guaranteed that $x^{(k)} \in C$ for some (finite) k .

Another method is to do *over-projection* at each step. Suppose we know the intersection of the sets contains a Euclidean ball of radius ϵ . Its center is a point that is ϵ -deep in all the sets. Then we can over project by ϵ , which roughly speaking means we project the current point to the farthest set, and then keep moving a distance ϵ :

$$x^{(k+1)} = \Pi_{C_j}(x^{(k)}) - \epsilon \frac{x^{(k)} - \Pi_{C_j}(x^{(k)})}{\|x^{(k)} - \Pi_{C_j}(x^{(k)})\|_2}.$$

Alternating projections is usually (but not always) applied when projection onto the sets is simple. This is the case, for example, for the following sets.

- Affine set.
- Nonnegative orthant.
- A halfspace or slab.
- A box, *e.g.*, unit ball in ℓ_∞ .
- Unit simplex.
- A Euclidean ball.
- An ellipsoid (there is no closed-form expression for the projection, but it can be computed very quickly.)
- A second-order cone.
- Cone of positive semidefinite matrices.
- Spectral norm matrix ball.

Alternating projections can be used, of course, in cases where a bit more computation is needed to compute the Euclidean projection, *e.g.*, for a polyhedron (which can be done by solving a QP).

5.2 Solving convex inequalities

We want to find a point that satisfies $f_i(x) \leq 0$, $i = 1, \dots, m$. (We assume we can find a subgradient of each function, at any point.)

To solve this set of convex inequalities, we can minimize the unconstrained function $f(x) = \max_i f_i(x)$ using the subgradient method. If the set of inequalities is strictly feasible, then f^* is negative, and in a finite number of steps we'll find a point with $f(x) \leq 0$, *i.e.*, a feasible point.

We can also use the step size that uses knowledge of the optimal value, applied to the function

$$f(x) = \max\{f_1(x), \dots, f_m(x), -\epsilon\},$$

where $\epsilon > 0$ is a tolerance. Assuming there exists a point with $f_i(x) \leq -\epsilon$, we can use the step length

$$\alpha = \frac{f(x) + \epsilon}{\|g\|_2^2}. \quad (6)$$

We can give a simple interpretation of this step length, taking the case $\epsilon = 0$ for simplicity. Suppose the current point is x , and that $f_i(x) = f(x) > 0$, with $g \in \partial f_i(x)$. Let x^* be any point with $f_i(x^*) \leq 0$. Then we have

$$0 \geq f_i(x^*) \geq f_i(x) + g^T(x^* - x),$$

i.e., x^* is in the halfspace

$$\mathcal{H} = \{z \mid 0 \geq f_i(x) + g^T(z - x)\}.$$

The subgradient update at x , using Polyak's step length, is just projection of x onto the halfspace \mathcal{H} .

As an example we consider finding a point $x \in \mathbf{R}^n$ that satisfies a set of linear inequalities $a_i^T x \leq b_i$, $i = 1, \dots, m$. With $\epsilon = 0$, the subgradient method is very simple: at each step, we find the most violated inequality. Then we project the current point onto the set (halfspace) of points that satisfy this particular inequality:

$$x^{(k+1)} = x^{(k)} - \frac{a_i^T x - b_i}{\|a_i\|_2^2} a_i,$$

where i is the index of the most violated inequality at $x^{(k)}$.

We take a problem instance with $n = 100$ variables and $m = 1000$ inequalities, and randomly generate the data, making sure that the set of inequalities is feasible. We use the step size rule (6) with three different values of ϵ . Figure 5 shows the convergence for $\epsilon = 0$, $\epsilon = 0.01$, and $\epsilon = 0.1$. (We terminate the algorithm when we find a point that satisfies the inequalities.)

5.3 Positive semidefinite matrix completion

We use the subgradient method with step size (5) to solve the *positive semidefinite matrix completion problem* (see [BV04, exer. 4.47]). We briefly describe the problem. Suppose we have a matrix in \mathbf{S}^n with some of its entries (including all of its diagonal entries) fixed, and the others to be found. The goal is to find values for the other entries so that the (completed) matrix is positive semidefinite.

We use alternating projections onto the set of positive semidefinite matrices \mathbf{S}_+^n , and the set of matrices with the given fixed entries. (Projection is in the Frobenius norm sense.)




Figure 5: Convergence of the maximum violation for the linear feasibility problem, where we use the subgradient method with Polyak’s step size and three different values of tolerance ϵ .

The first projection can be found from the eigenvalue decomposition (see [BV04, §8.1.1]); for example, let $X = \sum_{i=1}^n \lambda_i q_i q_i^T$, then

$$\Pi(X) = \sum_{i=1}^n \max\{0, \lambda_i\} q_i q_i^T.$$

The second projection is straightforward: we simply take the given matrix and set its fixed entries back to the given fixed values. Thus, the algorithm will alternate between eigenvalue decomposition and truncation, and re-setting the fixed entries back to their required values.

As a specific example we consider a randomly generated problem with a 50×50 matrix that is missing about half of its entries. The sparsity pattern of our particular matrix is shown in figure 6. We initialize $X^{(1)}$ by taking the unknown entries to be 0.

To track convergence of the algorithm, we plot the Frobenius norm of the difference between the current matrix and its projection onto one of the sets, *i.e.*, $\|X^{(k+1)} - X^{(k)}\|_F$. In the case of the projection onto the set of positive semidefinite matrices, this value is the squareroot of the sum of the squares of the negative eigenvalues of $X^{(k)}$. In the case of the other projection, it is the squareroot of the sum of the squares of the adjustments made to the fixed entries of $X^{(k)}$. In each case, this distance gives an upper bound on the distance to the intersection of the two sets, *i.e.*, the distance to the nearest positive semidefinite completion. The plot is shown in figure 7. We can see that the unknown entries are converging to a positive semidefinite completion. By overprojecting onto \mathbf{S}_+^n , we could have found an actual positive semidefinite completion in a finite number of steps.




Figure 6: Sparsity pattern of the given matrix with blue entries corresponding to fixed values and white entries corresponding to missing values.




Figure 7: Convergence of the subgradient method for a matrix completion problem.

6 Projected subgradient method

One extension of the subgradient method is the *projected subgradient method*, which solves the constrained convex optimization problem

$$\begin{aligned} & \text{minimize} && f(x) \\ & \text{subject to} && x \in \mathcal{C}, \end{aligned}$$

where \mathcal{C} is a convex set. The projected subgradient method is given by

$$x^{(k+1)} = \Pi \left(x^{(k)} - \alpha_k g^{(k)} \right),$$

where Π is (Euclidean) projection on \mathcal{C} , and $g^{(k)}$ is any subgradient of f at $x^{(k)}$. The step size rules described before can be used here, with similar convergence results. Note that $x^{(k)} \in \mathcal{C}$, *i.e.*, $x^{(k)}$ is feasible.

The convergence proofs for the subgradient method are readily extended to handle the projected subgradient method. Let $z^{(k+1)} = x^{(k)} - \alpha_k g^{(k)}$, *i.e.*, a standard subgradient update, before the projection back onto \mathcal{C} . As in the subgradient method, we have

$$\begin{aligned} \|z^{(k+1)} - x^*\|_2^2 &= \|x^{(k)} - \alpha_k g^{(k)} - x^*\|_2^2 \\ &= \|x^{(k)} - x^*\|_2^2 - 2\alpha_k g^{(k)T}(x^{(k)} - x^*) + \alpha_k^2 \|g^{(k)}\|_2^2 \\ &\leq \|x^{(k)} - x^*\|_2^2 - 2\alpha_k (f(x^{(k)}) - f^*) + \alpha_k^2 \|g^{(k)}\|_2^2. \end{aligned}$$

Now we observe that

$$\|x^{(k+1)} - x^*\|_2 = \|\Pi(z^{(k+1)}) - x^*\|_2 \leq \|z^{(k+1)} - x^*\|_2,$$

i.e., when we project a point onto \mathcal{C} , we move closer to every point in \mathcal{C} , and in particular, any optimal point. Combining this with the inequality above we get

$$\|x^{(k+1)} - x^*\|_2^2 \leq \|x^{(k)} - x^*\|_2^2 - 2\alpha_k (f(x^{(k)}) - f^*) + \alpha_k^2 \|g^{(k)}\|_2^2,$$

and the proof proceeds exactly as in the ordinary subgradient method. Also, note that Polyak's step size can be applied here directly, and has the same convergence guarantee.

In some cases we can express the projected subgradient update in an alternative way. When \mathcal{C} is affine, *i.e.*, $\mathcal{C} = \{x \mid Ax = b\}$, where A is fat and full rank, the projection operator is affine, and given by

$$\Pi(z) = z - A^T(AA^T)^{-1}(Az - b).$$

In this case, we can simplify the subgradient update to

$$x^{(k+1)} = x^{(k)} - \alpha_k (I - A^T(AA^T)^{-1}A)g^{(k)}, \tag{7}$$

where we use $Ax^{(k)} = b$. Thus, we simply project the current subgradient onto the nullspace of A , and then update as usual. The update (7) is not the same as the projected subgradient update when \mathcal{C} is not affine, because in this case the projection operator is not affine.




Figure 8: The value of $f_{\text{best}}^{(k)} - f^*$ versus iteration number k , for the subgradient method with the Polyak estimated step size rule $\gamma_k = 100/k$.

6.1 Numerical example

We consider the least l_1 -norm problem

$$\begin{aligned} & \text{minimize} && \|x\|_1 \\ & \text{subject to} && Ax = b, \end{aligned} \tag{8}$$

where the variable is $x \in \mathbf{R}^n$, and the data are $A \in \mathbf{R}^{m \times n}$ and $b \in \mathbf{R}^m$. We assume that A is fat and full rank, *i.e.*, $m < n$ and $\mathbf{Rank} A = m$. Of course, this problem is readily solved using linear programming.

A subgradient of the objective at x is given by $g = \mathbf{sign}(x)$. Thus, the projected subgradient update is

$$x^{(k+1)} = x^{(k)} - \alpha_k (I - A^T(AA^T)^{-1}A) \mathbf{sign}(x^{(k)}).$$

We consider an instance of the problem (8) with $n = 1000$ and $m = 50$, with randomly generated A and b . We use the least-norm solution as the starting point, *i.e.*, $x^{(1)} = A^T(AA^T)^{-1}b$. In order to report $f_{\text{best}}^{(k)} - f^*$, we solve the problem using linear programming and obtain $f^* \approx 3.2$. Figure 8 shows the progress of the projected subgradient method with the Polyak estimated step size rule $\gamma_k = 100/k$.

6.2 Projected subgradient for dual problem

One famous application of the projected subgradient method is to the dual problem. We start with the (convex) primal problem

$$\begin{aligned} & \text{minimize} && f_0(x) \\ & \text{subject to} && f_i(x) \leq 0, \quad i = 1, \dots, m \end{aligned}$$

We'll assume, for simplicity, that for each $\lambda \succeq 0$, the Lagrangian

$$L(x, \lambda) = f_0(x) + \sum_{i=1}^m \lambda_i f_i(x)$$

has a unique minimizer over x , which we denote $x^*(\lambda)$. The dual function is then

$$g(\lambda) = \inf_x L(x, \lambda) = f_0(x^*(\lambda)) + \sum_{i=1}^m \lambda_i f_i(x^*(\lambda))$$

(for $\lambda \succeq 0$). The dual problem is

$$\begin{aligned} & \text{maximize} && g(\lambda) \\ & \text{subject to} && \lambda \succeq 0. \end{aligned}$$

We'll assume that Slater's condition holds (again, for simplicity), so we can solve the primal problem by finding an optimal point λ^* of the dual, and then taking $x^* = x^*(\lambda^*)$. (For a discussion of solving the primal problem via the dual, see [BV04, §5.5.5].)

We will solve the dual problem using the projected subgradient method,

$$\lambda^{(k+1)} = \left(\lambda^{(k)} - \alpha_k h \right)_+, \quad h \in \partial(-g)(\lambda^{(k)}).$$

Let's now work out a subgradient of the negative dual function. Since $-g$ is a supremum of a family of affine functions of λ , indexed by x , we can find a subgradient by finding one of these functions that achieves the supremum. But there is just one, and it is

$$-f_0(x^*(\lambda)) - \sum_{i=1}^m \lambda_i f_i(x^*(\lambda)),$$

which has gradient (with respect to λ)

$$h = -(f_1(x^*(\lambda)), \dots, f_m(x^*(\lambda))) \in \partial(-g)(\lambda).$$

(Our assumptions imply that $-g$ has only one element in its subdifferential, which means g is differentiable. Differentiability means that a small enough constant step size will yield convergence. In any case, the projected subgradient method can be used in cases where the dual is nondifferentiable.)

The projected subgradient method for the dual has the form

$$\begin{aligned} x^{(k)} &= \operatorname{argmin}_x \left(f_0(x) + \sum_{i=1}^m \lambda_i^{(k)} f_i(x) \right) \\ \lambda_i^{(k+1)} &= \left(\lambda_i^{(k)} + \alpha_k f_i(x^{(k)}) \right)_+. \end{aligned}$$

In this algorithm, the primal iterates $x^{(k)}$ are not feasible, but become feasible only in the limit. (Sometimes we can find a method for constructing a feasible, suboptimal $\tilde{x}^{(k)}$ from $x^{(k)}$.) The dual function values $g(\lambda^{(k)})$, as well as the primal function values $f_0(x^{(k)})$, converge to $f^* = f_0(x^*)$.

We can interpret λ_i as the price for a ‘resource’ with usage measured by $f_i(x)$. When we calculate $x^*(\lambda)$, we are finding the x that minimizes the total cost, *i.e.*, the objective plus the total bill (or revenue) for the resources used. The goal is to adjust the prices so that the resource usage is within budget (*i.e.*, $f_i(x) \leq 0$). At each step, we increase the price λ_i if resource i is over-utilized (*i.e.*, $f_i(x) > 0$), and we decrease the price λ_i if resource i is under-utilized (*i.e.*, $f_i(x) < 0$). But we never let prices get negative (which would encourage, rather than discourage, resource usage).

In general, there is no reason to solve the dual instead of the primal. But for specific problems there can be an advantage.

6.3 Numerical example

We consider the problem of minimizing a strictly convex quadratic function over the unit box:

$$\begin{aligned} &\text{minimize} && (1/2)x^T P x - q^T x \\ &\text{subject to} && x_i^2 \leq 1, \quad i = 1, \dots, n, \end{aligned}$$

where $P \succ 0$. The Lagrangian is

$$L(x, \lambda) = (1/2)x^T (P + \mathbf{diag}(2\lambda))x - q^T x - \mathbf{1}^T \lambda,$$

so $x^*(\lambda) = (P + \mathbf{diag}(2\lambda))^{-1}q$. The projected subgradient algorithm for the dual is

$$x^{(k)} = (P + \mathbf{diag}(2\lambda^{(k)}))^{-1}q, \quad \lambda_i^{(k+1)} = \left(\lambda_i^{(k)} + \alpha_k ((x_i^{(k)})^2 - 1) \right)_+.$$

The dual function is differentiable, so we can use a fixed size α (provided it is small enough).

The iterates $x^{(k)}$ are not feasible. But we can construct a nearby feasible $\tilde{x}^{(k)}$ as

$$\tilde{x}_i^{(k)} = \begin{cases} 1 & x_i^{(k)} > 1 \\ -1 & x_i^{(k)} < -1 \\ x_i^{(k)} & -1 \leq x_i^{(k)} \leq 1. \end{cases}$$

We consider an instance with $n = 50$. We start the algorithm with $\lambda^{(1)} = \mathbf{1}$, and use a fixed step size $\alpha = 0.1$. Figure 9 shows the convergence of $g(\lambda^{(k)})$ (a lower bound on the optimal value) and $f_0(\tilde{x}^{(k)})$ (an upper bound on the optimal value), versus iterations.




Figure 9: The values of the lower bound $g(\lambda^{(k)})$ and the upper bound $f_0(\tilde{x}^{(k)})$, versus the iteration number k . We use the fixed step size with $\alpha = 0.1$.

7 Subgradient method for constrained optimization

The subgradient algorithm can be extended to solve the inequality constrained problem

$$\begin{aligned} & \text{minimize} && f_0(x) \\ & \text{subject to} && f_i(x) \leq 0, \quad i = 1, \dots, m, \end{aligned}$$

where f_i are convex. The algorithm takes the same form:

$$x^{(k+1)} = x^{(k)} - \alpha_k g^{(k)},$$

where $\alpha_k > 0$ is a step size, and $g^{(k)}$ is a subgradient of the objective or one of the constraint functions at $x^{(k)}$. More specifically, we take

$$g^{(k)} \in \begin{cases} \partial f_0(x^{(k)}) & f_i(x^{(k)}) \leq 0, \quad i = 1, \dots, m, \\ \partial f_j(x^{(k)}) & f_j(x^{(k)}) > 0. \end{cases}$$

In other words: If the current point is feasible, we use an objective subgradient, as if the problem were unconstrained; if the current point is infeasible, we choose any violated constraint, and use a subgradient of the associated constraint function. (In the latter case, we can choose *any* of the violated constraints, if there is more than one.)

In this generalized version of the subgradient algorithm, the iterates can be (and often are) infeasible. In contrast, the iterates of the projected subgradient method (and of course, the basic subgradient algorithm) are always feasible.

As in the basic subgradient method, we keep track of the best (feasible) point found so far:

$$f_{\text{best}}^{(k)} = \min\{f_0(x^{(i)}) \mid x^{(i)} \text{ feasible}, i = 1, \dots, k\}.$$

(If none of the points $x^{(1)}, \dots, x^{(k)}$ is feasible, then $f_{\text{best}}^{(k)} = \infty$.)

We assume that Slater's condition holds, *i.e.*, the problem is strictly feasible: there is some point x^{sf} with $f_i(x^{\text{sf}}) < 0$, $i = 1, \dots, m$. We also assume that the problem has an optimal point x^* . We assume that there are numbers R and G with $\|x^{(1)} - x^*\|_2 \leq R$, $\|x^{(1)} - x^{\text{sf}}\|_2 \leq R$, and $\|g^{(k)}\|_2 \leq G$ for all k .

We'll establish convergence of the generalized subgradient method using diminishing non-summable α_k . (Similar results can be obtained for other step size rules.) We claim that $f_{\text{best}}^{(k)} \rightarrow f^*$ as $k \rightarrow \infty$. This implies in particular that we obtain a feasible iterate within some finite number of steps.

Assume that $f_{\text{best}}^{(k)} \rightarrow f^*$ does not occur. Then there exists some $\epsilon > 0$ so that $f_{\text{best}}^{(k)} \geq f^* + \epsilon$ for all k , which in turn means that $f(x^{(k)}) \geq f^* + \epsilon$ for all k for which $x^{(k)}$ is feasible. We'll show this leads to a contradiction.

We first find a point \tilde{x} and positive number μ that satisfy

$$f_0(\tilde{x}) \leq f^* + \epsilon/2, \quad f_1(\tilde{x}) \leq -\mu, \dots, f_m(\tilde{x}) \leq -\mu.$$

Such a point is $\epsilon/2$ -suboptimal, and also satisfies the constraints with a margin of μ . We will take $\tilde{x} = (1 - \theta)x^* + \theta x^{\text{sf}}$, where $\theta \in (0, 1)$. We have

$$f_0(\tilde{x}) \leq (1 - \theta)f^* + \theta f_0(x^{\text{sf}}),$$

so if we choose $\theta = \min\{1, (\epsilon/2)/(f_0(x^{\text{sf}}) - f^*)\}$, we have $f_0(\tilde{x}) \leq f^* + \epsilon/2$. We have

$$f_i(\tilde{x}) \leq (1 - \theta)f_i(x^*) + \theta f_i(x^{\text{sf}}) \leq \theta f_i(x^{\text{sf}}),$$

so we can take

$$\mu = -\theta \min_i f_i(x^{\text{sf}}).$$

Consider any index $i \in \{1, \dots, k\}$ for which $x^{(i)}$ is feasible. Then we have $g^{(i)} \in \partial f_0(x^{(i)})$, and also $f_0(x^{(i)}) \geq f^* + \epsilon$. Since \tilde{x} is $\epsilon/2$ -suboptimal, we have $f_0(x^{(i)}) - f_0(\tilde{x}) \geq \epsilon/2$. Therefore

$$\begin{aligned} \|x^{(i+1)} - \tilde{x}\|_2^2 &= \|x^{(i)} - \tilde{x}\|_2^2 - 2\alpha_i g^{(i)T}(x^{(i)} - \tilde{x}) + \alpha_i^2 \|g^{(i)}\|_2^2 \\ &\leq \|x^{(i)} - \tilde{x}\|_2^2 - 2\alpha_i (f_0(x^{(i)}) - f_0(\tilde{x})) + \alpha_i^2 \|g^{(i)}\|_2^2 \\ &\leq \|x^{(i)} - \tilde{x}\|_2^2 - \alpha_i \epsilon + \alpha_i^2 \|g^{(i)}\|_2^2. \end{aligned}$$

In the second line here we use the usual subgradient inequality

$$f_0(\tilde{x}) \geq f_0(x^{(i)}) + g^{(i)T}(\tilde{x} - x^{(i)}).$$

Now suppose that $i \in \{1, \dots, k\}$ is such that $x^{(i)}$ is infeasible, and that $g^{(i)} \in \partial f_p(x^{(i)})$, where $f_p(x^{(i)}) > 0$. Since $f_p(\tilde{x}) \leq -\mu$, we have $f_p(x^{(i)}) - f_p(\tilde{x}) \geq \mu$. Therefore

$$\begin{aligned} \|x^{(i+1)} - \tilde{x}\|_2^2 &= \|x^{(i)} - \tilde{x}\|_2^2 - 2\alpha_i g^{(i)T}(x^{(i)} - \tilde{x}) + \alpha_i^2 \|g^{(i)}\|_2^2 \\ &\leq \|x^{(i)} - \tilde{x}\|_2^2 - 2\alpha_i (f_p(x^{(i)}) - f_p(\tilde{x})) + \alpha_i^2 \|g^{(i)}\|_2^2 \\ &\leq \|x^{(i)} - \tilde{x}\|_2^2 - 2\alpha_i \mu + \alpha_i^2 \|g^{(i)}\|_2^2. \end{aligned}$$

Thus, for every iteration we have

$$\|x^{(i+1)} - \tilde{x}\|_2^2 \leq \|x^{(i)} - \tilde{x}\|_2^2 - \alpha_i \delta + \alpha_i^2 \|g^{(i)}\|_2^2,$$

where $\delta = \min\{\epsilon, 2\mu\} > 0$. Applying this inequality recursively for $i = 1, \dots, k$, we get

$$\|x^{(k+1)} - \tilde{x}\|_2^2 \leq \|x^{(1)} - \tilde{x}\|_2^2 - \delta \sum_{i=1}^k \alpha_i + \sum_{i=1}^k \alpha_i^2 \|g^{(i)}\|_2^2.$$

It follows that

$$\delta \sum_{i=1}^k \alpha_i \leq R^2 + G^2 \sum_{i=1}^k \alpha_i^2,$$

which cannot hold for large k since

$$\frac{R^2 + G^2 \sum_{i=1}^k \alpha_i^2}{\sum_{i=1}^k \alpha_i}$$

converges to zero as $k \rightarrow \infty$.

There are many variations on the basic step size rule. For example, when the current point is infeasible, we can use an over-projection step length, as we would when solving

convex inequalities. If we know (or estimate) f^\star , we can use Polyak's step length when the current point is feasible. Thus our step lengths are chosen as

$$\alpha_k = \begin{cases} (f_0(x^{(k)}) - f^\star) / \|g^{(k)}\|_2^2 & x^{(k)} \text{ feasible} \\ (f_i(x^{(k)}) + \epsilon) / \|g^{(k)}\|_2^2 & x^{(k)} \text{ infeasible} \end{cases}$$

where ϵ is a small positive margin, and i is the index of the most violated inequality in the case when $x^{(k)}$ is infeasible.

7.1 Numerical example

We consider a linear program

$$\begin{aligned} & \text{minimize} && c^T x \\ & \text{subject to} && a_i^T x \leq b_i, \quad i = 1, \dots, m, \end{aligned} \tag{9}$$

with variable $x \in \mathbf{R}^n$. The objective and constraint functions are affine, and so have only one subgradient, independent of x . For the objective function we have $g = c$, and for the i th constraint we have $g_i = a_i$.

We solve an instance of the problem (9) with $n = 20$ and $m = 200$ using the subgradient method. In order to report $f_{\text{best}}^{(k)} - f^\star$, we solve the LP using the interior-point methods and obtain $f^\star \approx -3.4$. Figure 10 shows progress of the subgradient method, which uses the square summable step size with $\alpha_k = 1/k$ for the optimality update, and the step size (6) with $\epsilon = 10^{-3}$ for the feasibility update. The objective value only changes for the iterations when $x^{(k)}$ is feasible.




Figure 10: The value of $f_{\text{best}}^{(k)} - f^*$ versus the iteration number k . In this case, we use the square summable step size with $\alpha_k = 1/k$ for the optimality update.

8 Primal-dual subgradient method

The *primal-dual subgradient method* is an extension of the subgradient method that solves constrained convex optimization problem of the form

$$\begin{aligned} & \text{minimize} && f_0(x) \\ & \text{subject to} && f_i(x) \leq 0, \quad i = 1, \dots, m \\ & && Ax = b, \end{aligned}$$

with variable $x \in \mathbf{R}^n$, where the functions $f_0, \dots, f_m : \mathbf{R}^n \rightarrow \mathbf{R}$ are convex, not necessarily differentiable, and have domain \mathbf{R}^n . In the following sections, we will consider problems that have the equality constraint only, or inequality constraints only, and show the method with proof of its convergence. Extending these ideas to handle problems with both inequality and equality constraints is simple.

8.1 Equality constrained problems

In this section, we consider the equality constrained problem

$$\begin{aligned} & \text{minimize} && f(x) \\ & \text{subject to} && Ax = b, \end{aligned}$$

with variable $x \in \mathbf{R}^n$. We assume that $A \in \mathbf{R}^{m \times n}$, *i.e.*, there are m equality constraints.

Note that $(\rho/2)\|Ax - b\|_2^2$, where $\rho > 0$, can be added to the objective without changing the optimal value or the set of optimal points. Thus, from now on, we focus on solving the so-called *augmented problem*

$$\begin{aligned} & \text{minimize} && f(x) + (\rho/2)\|Ax - b\|_2^2 \\ & \text{subject to} && Ax = b. \end{aligned}$$

Let $L(x, \nu) = f(x) + \nu^T(Ax - b) + (\rho/2)\|Ax - b\|_2^2$ denote the Lagrangian for the augmented problem, which is also called the *augmented Lagrangian*. We define a set-valued mapping T by

$$T(x, \nu) = \begin{bmatrix} \partial_x L(x, \nu) \\ -\partial_\nu L(x, \nu) \end{bmatrix} = \begin{bmatrix} \partial f(x) + A^T \nu + \rho A^T (Ax - b) \\ b - Ax \end{bmatrix}.$$

(We mention that T is the KKT operator associated with the augmented Lagrangian, and is a monotone operator.) The optimality condition for the augmented problem (and the original one as well) is

$$0 \in T(x^*, \nu^*).$$

Such a primal-dual pair is a saddle-point of the augmented Lagrangian:

$$L(x^*, \nu) \leq L(x^*, \nu^*) \leq L(x, \nu^*)$$

for all x and all ν .

The primal-dual subgradient method finds a saddle point of the Lagrangian by the simple iteration that resembles the Uzawa iteration (see [BPV97]):

$$z^{(k+1)} = z^{(k)} - \alpha_k T^{(k)},$$

where $z^{(k)} = (x^{(k)}, \nu^{(k)})$ is the k th iterate of the primal and dual variables, $T^{(k)}$ is any element of $T(z^{(k)})$, and $\alpha_k > 0$ is the k th step size. By expanding it out, we can also write the method as

$$\begin{aligned} x^{(k+1)} &= x^{(k)} - \alpha_k (g^{(k)} + A^T \nu^{(k)} + \rho A^T (Ax^{(k)} - b)) \\ \nu^{(k+1)} &= \nu^{(k)} + \alpha_k (Ax^{(k)} - b). \end{aligned}$$

Here, $g^{(k)}$ is any subgradient of f at $x^{(k)}$. Notice that $x^{(k)}$ is not necessarily feasible.

Let $z^* = (x^*, \nu^*)$ be any pair of optimal variables, satisfying

$$Ax^* = b, \quad 0 \in \partial_x L(x^*, \nu^*).$$

We use $p^* = f(x^*)$ to denote the optimal value. We prove that the algorithm converges, *i.e.*,

$$\lim_{k \rightarrow \infty} f(x^{(k)}) = p^*, \quad \lim_{k \rightarrow \infty} \|Ax^{(k)} - b\|_2 = 0,$$

using the step size rule $\alpha_k = \gamma_k / \|T^{(k)}\|_2$, where γ_k is chosen so that

$$\gamma_k > 0, \quad \sum_{k=1}^{\infty} \gamma_k = \infty, \quad \sum_{k=1}^{\infty} \gamma_k^2 = S < \infty.$$

For the convergence proof, we will assume that a number R is known that satisfies $R \geq \|z^{(1)}\|_2$ and $R \geq \|z^*\|_2$. We will also assume that the norm of the subgradients of f is bounded on compact sets.

We start by writing out a basic identity

$$\begin{aligned} \|z^{(k+1)} - z^*\|_2^2 &= \|z^{(k)} - z^*\|_2^2 - 2\alpha_k T^{(k)T}(z^{(k)} - z^*) + \alpha_k^2 \|T^{(k)}\|_2^2 \\ &= \|z^{(k)} - z^*\|_2^2 - 2\gamma_k \frac{T^{(k)T}}{\|T^{(k)}\|_2} (z^{(k)} - z^*) + \gamma_k^2. \end{aligned}$$

By summing it over k and rearranging the terms, we get

$$\|z^{(k+1)} - z^*\|_2^2 + 2 \sum_{i=0}^k \gamma_i \frac{T^{(i)T}}{\|T^{(i)}\|_2} (z^{(i)} - z^*) = \|z^{(1)} - z^*\|_2^2 + \sum_{i=0}^k \gamma_i^2 \leq 4R^2 + S. \quad (10)$$

We argue that the sum on the lefthand side is nonnegative. The monotonicity of T directly implies this, but it can be shown without using monotonicity. First, we expand out as

$$T^{(k)T}(z^{(k)} - z^*) = (g^{(k)} + A^T \nu^{(k)} + \rho A^T (Ax^{(k)} - b))^T (x^{(k)} - x^*) - (Ax^{(k)} - b)^T (\nu^{(k)} - \nu^*).$$

Since $Ax^\star = b$, the first term further expands to

$$g^{(k)T}(x^{(k)} - x^\star) + \nu^{(k)T}(Ax^{(k)} - b) + \rho \|Ax^{(k)} - b\|_2^2.$$

By subtracting the second term and using the definition of subgradient,

$$\begin{aligned} T^{(k)T}(z^{(k)} - z^\star) &= g^{(k)T}(x^{(k)} - x^\star) + \nu^{\star T}(Ax^{(k)} - b) + \rho \|Ax^{(k)} - b\|_2^2 \\ &\geq f(x^{(k)}) - p^\star + \nu^{\star T}(Ax^{(k)} - b) + \rho \|Ax^{(k)} - b\|_2^2 \\ &= L(x^{(k)}, \nu^\star) - L(x^\star, \nu^\star) + (\rho/2) \|Ax^{(k)} - b\|_2^2 \\ &\geq 0. \end{aligned}$$

The last line uses the fact that x^\star minimizes $L(x, \nu^\star)$ over x .

Since both terms on the lefthand side of (10) are nonnegative, for all k , we have

$$\|z^{(k+1)} - z^\star\|_2^2 \leq 4R^2 + S, \quad 2 \sum_{i=0}^k \gamma_i \frac{T^{(i)T}}{\|T^{(i)}\|_2} (z^{(i)} - z^\star) \leq 4R^2 + S,$$

for all k . We assumed that $\|z^\star\|_2$ is bounded, so the first inequality implies that $z^{(k)}$ cannot be too far from the origin. In other words, there exists a number D satisfying $\|z^{(k)}\|_2 \leq D$ for all k , namely $D = R + \sqrt{4R^2 + S}$. By assumption, the norm of subgradients on the set $\|x^{(k)}\|_2 \leq D$ is bounded, so it follows that $\|T^{(k)}\|_2$ is bounded.

Because the sum of γ_k diverges, for the sum

$$\sum_{i=0}^k \gamma_i \frac{T^{(i)T}}{\|T^{(i)}\|_2} (z^{(i)} - z^\star)$$

to be bounded, we need

$$\lim_{k \rightarrow \infty} \frac{T^{(k)T}}{\|T^{(k)}\|_2} (z^{(k)} - z^\star) = 0.$$

However, since $\|T^{(k)}\|_2$ is bounded, the numerator $T^{(k)T}(z^{(k)} - z^\star)$ has to go to zero in the limit. Note that the inequality

$$0 \leq L(x^{(k)}, \nu^\star) - L(x^\star, \nu^\star) + (\rho/2) \|Ax^{(k)} - b\|_2^2 \leq T^{(k)T}(z^{(k)} - z^\star),$$

together with the fact that $L(x^{(k)}, \nu^\star) - L(x^\star, \nu^\star) \geq 0$ and $(\rho/2) \|Ax^{(k)} - b\|_2^2 \geq 0$, implies that

$$\lim_{k \rightarrow \infty} L(x^{(k)}, \nu^\star) = L(x^\star, \nu^\star) = p^\star, \quad \lim_{k \rightarrow \infty} \|Ax^{(k)} - b\|_2 = 0.$$

Finally,

$$p^\star = \lim_{k \rightarrow \infty} L(x^{(k)}, \nu^\star) = \lim_{k \rightarrow \infty} f(x^{(k)}) + \lim_{k \rightarrow \infty} \nu^{\star T}(Ax^{(k)} - b) = \lim_{k \rightarrow \infty} f(x^{(k)}).$$

This proves the convergence of the algorithm.

8.2 Inequality constrained problems

The method from the previous section can be modified to handle inequality constrained problems. Suppose that we want to solve the problem

$$\begin{aligned} & \text{minimize} && f_0(x) \\ & \text{subject to} && f_i(x) \leq 0, \quad i = 1, \dots, m. \end{aligned}$$

with variable $x \in \mathbf{R}^n$.

Notice that $(\rho/2) \sum_{i=1}^m (f_i(x)_+)^2$, with $\rho > 0$, can be added to the objective without changing the optimal value or the set of optimal. Thus, from now on, we focus on solving the augmented problem

$$\begin{aligned} & \text{minimize} && f_0(x) + (\rho/2) \|F(x)\|_2^2 \\ & \text{subject to} && F(x) \preceq 0, \end{aligned}$$

where F is defined by

$$F(x) = \begin{bmatrix} f_1(x)_+ \\ \vdots \\ f_m(x)_+ \end{bmatrix}.$$

Let $L(x, \lambda) = f_0(x) + \lambda^T F(x) + (\rho/2) \|F(x)\|_2^2$ be the augmented Lagrangian. We define a set-valued mapping T by

$$T(x, \lambda) = \begin{bmatrix} \partial_x L(x, \lambda) \\ -\partial_\lambda L(x, \lambda) \end{bmatrix} = \begin{bmatrix} \partial f_0(x) + \sum_{i=1}^m (\lambda_i + \rho f_i(x)_+) \partial f_i(x)_+ \\ -F(x) \end{bmatrix}.$$

The optimality condition for the augmented problem is then

$$0 \in T(x^*, \lambda^*).$$

Such a primal-dual pair is a saddle-point of the augmented Lagrangian:

$$L(x^*, \lambda) \leq L(x^*, \lambda^*) \leq L(x, \lambda^*)$$

for all x and all λ .

The primal-dual subgradient method can be written as

$$z^{(k+1)} = z^{(k)} - \alpha_k T^{(k)},$$

where $z^{(k)} = (x^{(k)}, \lambda^{(k)})$ is the k th iterate of the primal and dual variables, $T^{(k)}$ is any element of $T(z^{(k)})$, and $\alpha_k > 0$ is the k th step size. By expanding it out, we can also write the method as

$$\begin{aligned} x^{(k+1)} &= x^{(k)} - \alpha_k \left(g_0^{(k)} + \sum_{i=1}^m (\lambda_i^{(k)} + \rho f_i(x^{(k)})_+) g_i^{(k)} \right) \\ \lambda_i^{(k+1)} &= \lambda_i^{(k)} + \alpha_k f_i(x^{(k)}), \quad i = 1, \dots, m. \end{aligned}$$

Here, $g_i^{(k)}$ is any subgradient of f_i at $x^{(k)}$.

The convergence proof goes similarly. Let $z^* = (x^*, \lambda^*)$ be any pair of optimal variables. That is,

$$F(x^*) = 0, \quad 0 \in \partial_x L(x^*, \lambda^*).$$

Let $p^* = f_0(x^*)$ denote the optimal value. We prove that

$$\lim_{k \rightarrow \infty} f_0(x^{(k)}) = p^*, \quad \lim_{k \rightarrow \infty} \|F(x^{(k)})\|_2 = 0,$$

using the step size rule $\alpha_k = \gamma_k / \|T^{(k)}\|_2$, where γ_k is square summable but not summable. We will also assume that the norm of the subgradients of each f_i , and the values of f_1, \dots, f_m are bounded on compact sets.

We start with the same basic inequality obtained in the previous section:

$$\|z^{(k+1)} - z^*\|_2^2 + 2 \sum_{i=0}^k \gamma_i \frac{T^{(i)T}}{\|T^{(i)}\|_2} (z^{(i)} - z^*) \leq 4R^2 + S.$$

Again, we claim that the sum on the lefthand side is nonnegative. By expanding out,

$$\begin{aligned} T^{(k)T}(z^{(k)} - z^*) &= \left(g_0^{(k)} + \sum_{i=1}^m (\lambda_i^{(k)} + \rho f_i(x^{(k)})_+) g_i^{(k)} \right)^T (x^{(k)} - x^*) \\ &\quad - F(x^{(k)})^T (\lambda^{(k)} - \lambda^*). \end{aligned}$$

Here, $g_i^{(k)}$ denotes any subgradient of $(f_i)_+$ at $x^{(k)}$. By definition of subgradient, for the objective function, we have

$$g_0^{(k)T}(x^{(k)} - x^*) \geq f_0(x^{(k)}) - p^*,$$

and for the constraints,

$$g_i^{(k)T}(x^{(k)} - x^*) \geq f_i(x^{(k)})_+ - f_i(x^*)_+ = f_i(x^{(k)})_+, \quad i = 1, \dots, m.$$

Using these, we have

$$\begin{aligned} T^{(k)T}(z^{(k)} - z^*) &\geq f_0(x^{(k)}) - p^* + \lambda^{*T} F(x^{(k)}) + \rho \|F(x^{(k)})\|_2^2 \\ &= L(x^{(k)}, \lambda^*) - L(x^*, \lambda^*) + (\rho/2) \|F(x^{(k)})\|_2^2 \\ &\geq 0. \end{aligned}$$

The last line is true because x^* minimizes $L(x, \lambda^*)$ over x . The rest of the proof proceeds exactly the same as in the case of equality constrained problems.

8.3 Numerical example

We consider a linear program

$$\begin{aligned} & \text{minimize} && c^T x \\ & \text{subject to} && a_i^T x \leq b_i, \quad i = 1, \dots, m, \end{aligned}$$

with variable $x \in \mathbf{R}^n$. The objective is affine, so the only subgradient of it, independent of x , is given by $g_0 = c$. For the i th constraint, the following is a valid subgradient at x :

$$g_i = \begin{cases} a_i & a_i^T x > b_i \\ 0 & a_i^T x \leq b_i. \end{cases}$$

Let A be the matrix whose i th row is a_i^T . We can explicitly write the update rule for the primal-dual method as

$$\begin{aligned} x^{(k+1)} &= x^{(k)} - \alpha_k \left(c + A^T M^{(k)} (\lambda^{(k)} + \rho(Ax^{(k)} - b)_+) \right) \\ \lambda^{(k+1)} &= \lambda^{(k)} + \alpha_k (Ax^{(k)} - b)_+, \end{aligned}$$

where the notation x_+ denotes the vector obtained by taking the positive part of each element of x , and $M^{(k)}$ is a diagonal matrix given by

$$M_{ii}^{(k)} = \begin{cases} 1 & a_i^T x^{(k)} > b_i \\ 0 & a_i^T x^{(k)} \leq b_i. \end{cases}$$

We solve the same instance of the problem appeared in §7.1 using the primal-dual subgradient method. In order to report $f_{\text{best}}^{(k)} - p^*$, we solve the LP using the interior-point methods and obtain $p^* \approx -3.4$. Figure 11 shows progress of the primal-dual subgradient method, which uses the step size $\gamma_k = 1/k$.




Figure 11: The suboptimality $|f(x^{(k)}) - p^*|$ (in blue), and the maximum violation of the constraints $\max_{i=1,\dots,m}(a_i^T x - b_i)$ (in red), versus the iteration number k . In this case, we use the square summable sequence $\gamma_k = 1/k$ to determine the step sizes.

9 Speeding up subgradient methods

Several general approaches can be used to speed up subgradient methods. *Localization methods* such as cutting-plane and ellipsoid methods also require the evaluation of one subgradient per iteration, but require more computation to carry out the update. They are typically much faster than subgradient methods. Some of these methods have real (non-heuristic) stopping criteria.

9.1 Heavy ball methods

Another general approach is to base the update on some conic combination of previously evaluated subgradients. In *bundle methods*, the update direction is found as the least-norm convex combination of some ('bundle' of) previous subgradients. (This gives an approximation of the steepest descent direction.)

One general class of methods uses an update direction that is a conic combination of the current negative subgradient and the last search direction, as in

$$x^{(k+1)} = x^{(k)} - \alpha_k g^{(k)} + \beta_k (x^{(k)} - x^{(k-1)})$$

where α_k and β_k are positive. (There are many other ways to express this update.) Such algorithms have state, whereas the basic subgradient method is stateless (except for the iteration number). We can interpret the second term as a memory term, or as a momentum term, in the algorithm. Polyak refers to some algorithms of this form as the *heavy ball method*. *Conjugate gradients* methods have a similar form.

We describe two examples of these types of methods, that use a known (or estimated) value of f^* to determine step lengths. Each has an update of the form

$$x^{(k+1)} = x^{(k)} - \alpha_k s^{(k)}, \quad \alpha_k = \frac{f(x^{(k)}) - f^*}{\|s^{(k)}\|_2^2},$$

where $s^{(k)}$ is a direction to be used in place of a subgradient. In the simple method, $s^{(k)}$ is just a filtered, or smoothed, version of the subgradients:

$$s^{(k)} = (1 - \beta)g^{(k)} + \beta s^{(k-1)},$$

where $0 \leq \beta < 1$ is a (constant) filter parameter that controls how much memory the algorithm has. When $\beta = 0$ we obtain the subgradient method with Polyak's step size.

A more sophisticated method for updating $s^{(k)}$ was proposed by Camerini, Fratta, and Maffioli [CFM75]. Their algorithm has the form

$$s^{(k)} = g^{(k)} + \beta_k s^{(k-1)}, \tag{11}$$

where

$$\beta_k = \max\{0, -\gamma_k (s^{(k-1)})^T g^{(k)} / \|s^{(k-1)}\|_2^2\}.$$

Here $\gamma_k \in [0, 2]$; they recommend using the constant value $\gamma_k = 1.5$.




Figure 12: The value of $f_{\text{best}}^{(k)} - f^*$ versus iteration number k , for the subgradient method with two types of Polyak's step sizes, the original update when $\beta = 0$ (dashed black line) and a filtered update with $\beta = 0.25$ (solid blue line). The plot also shows the subgradient method with CFM step size (dash dotted green line).

They show that

$$\frac{(x^{(k)} - x^*)^T s^{(k)}}{\|s^{(k)}\|_2^2} \geq \frac{(x^{(k)} - x^*)^T g^{(k)}}{\|g^{(k)}\|_2^2},$$

i.e., the direction with modified update has a smaller angle towards the optimal set than the negative subgradient. (It follows that the convergence proofs for the subgradient algorithm work for this one as well.)

To illustrate these acceleration techniques, we consider again our piecewise-linear minimization example. We use the CFM algorithm and its simpler update rule given above for $\beta = 0$ and $\beta = 0.25$. Figure 12 shows the progress of these algorithms.

Acknowledgments

We thank Lieven Vandenberghe, Lin Xiao, and Almir Mutapcic, who helped with earlier versions of these notes.

References

- [Akg84] M. Akgül. *Topics in Relaxation and Ellipsoidal Methods*, volume 97 of *Research Notes in Mathematics*. Pitman, 1984.
- [Ber99] D. Bertsekas. *Nonlinear Programming*. Athena Scientific, second edition, 1999.
- [BPV97] J. Bramble, J. Pasciak, and A. Vassilev. Analysis of the inexact uzawa algorithm for saddle point problems. *SIAM Journal on Numerical Analysis*, 34(3):1072–1092, 1997.
- [BV04] S. Boyd and L. Vandenberghe. *Convex Optimization*. Cambridge University Press, 2004.
- [CFM75] P. Camerini, L. Fratta, and F. Maffioli. On improving relaxation methods by modifying gradient techniques. *Math. Programming Study*, 3:26–34, 1975.
- [CZ97] Y. Censor and S. Zenios. *Parallel Optimization*. Oxford University Press, 1997.
- [NB01] A. Nedić and D. Bertsekas. Incremental subgradient methods for nondifferentiable optimization. *SIAM J. on Optimization*, 12:109–138, 2001.
- [Nes04] Y. Nesterov. *Introductory Lectures on Convex Optimization: A Basic Course*. Kluwer Academic Publishers, 2004.
- [Nes09] Y. Nesterov. Primal-dual subgradient methods for convex problems. *Mathematical Programming*, 120(1):221–259, 2009.
- [NY83] A. Nemirovski and D. Yudin. *Problem Complexity and Method Efficiency in Optimization*. Wiley-Interscience, 1983.
- [Pol87] B. Polyak. *Introduction to Optimization*. Optimization Software, Inc., 1987.
- [Rus06] A. Ruszczyński. *Nonlinear Optimization*. Princeton University Press, 2006.
- [Sho85] N. Shor. *Minimization Methods for Non-differentiable Functions*. Springer Series in Computational Mathematics. Springer, 1985.
- [Sho98] N. Shor. *Nondifferentiable Optimization and Polynomial Problems*. Nonconvex Optimization and its Applications. Kluwer, 1998.